

# AIFS ML Lecture 1: Machine Learning Basics

Suraj Narayanan Sasikumar

Hessian AI Labs

April 15, 2020



- 1 Module Outline
- 2 Introduction
  - What is Intelligence?
  - What is Machine Learning?
- 3 Learning Algorithms
- 4 Reference



# Table of Contents

## 1 Module Outline

## 2 Introduction

- What is Intelligence?
- What is Machine Learning?

## 3 Learning Algorithms

## 4 Reference



# Module Outline

- Machine Learning Basics
- Linear Models for Regression and Classifications
- Non-linear Models
- Kernel Methods
- Graphical Models
- Latent Variable Models
- Approximate Inference and Sampling
- Dimensionality Reduction

## Textbooks:

- Pattern Recognition and Machine Learning Christopher M. Bishop
- Deep Learning, Ian Goodfellow et. al.(Ch5)



# Table of Contents

## 1 Module Outline

## 2 Introduction

- What is Intelligence?
- What is Machine Learning?

## 3 Learning Algorithms

## 4 Reference



# What is Intelligence?

- What is the goal of Artificial Intelligence (AI)?



# What is Intelligence?

- What is the goal of Artificial Intelligence (AI)?
- Simply put, the goal of AI is to enable machines to behave more intelligently.



# What is Intelligence?

- What is the goal of Artificial Intelligence (AI)?
- Simply put, the goal of AI is to enable machines to behave more intelligently.
- If it can behave as intelligently as humans, we call it an Artificial General Intelligence (AGI). If it can only do well in a restricted set of tasks then we call it an Artificial Narrow Intelligence (ANI)<sup>1</sup>.

---

<sup>1</sup>In this course we'll focus only on ANI.





# What is Intelligence?

- What is the goal of Artificial Intelligence (AI)?
- Simply put, the goal of AI is to enable machines to behave more intelligently.
- If it can behave as intelligently as humans, we call it an Artificial General Intelligence (AGI). If it can only do well in a restricted set of tasks then we call it an Artificial Narrow Intelligence (ANI)<sup>1</sup>.
- So what is intelligence, or how do you measure it?

---

<sup>1</sup>In this course we'll focus only on ANI.

# What is Intelligence?

- What is the goal of Artificial Intelligence (AI)?
- Simply put, the goal of AI is to enable machines to behave more intelligently.
- If it can behave as intelligently as humans, we call it an Artificial General Intelligence (AGI). If it can only do well in a restricted set of tasks then we call it an Artificial Narrow Intelligence (ANI)<sup>1</sup>.
- So what is intelligence, or how do you measure it?
- Intelligence is a measure of some agent's (here a computer program) ability to achieve its goals in a wide range of environments (Legg and Hutter, 2007).

---

<sup>1</sup>In this course we'll focus only on ANI.

# What is Intelligence?

- What is the goal of Artificial Intelligence (AI)?
- Simply put, the goal of AI is to enable machines to behave more intelligently.
- If it can behave as intelligently as humans, we call it an Artificial General Intelligence (AGI). If it can only do well in a restricted set of tasks then we call it an Artificial Narrow Intelligence (ANI)<sup>1</sup>.
- So what is intelligence, or how do you measure it?
- Intelligence is a measure of some agent's (here a computer program) ability to achieve its goals in a wide range of environments (Legg and Hutter, 2007).
- For example, we can say that a computer program that is able to play both Chess and Go (AlphaZero) is more intelligent than a program that can only play Go (AlphaGo).

---

<sup>1</sup>In this course we'll focus only on ANI.

# Prerequisites for Intelligence

- Suppose Alice is a layperson who doesn't know anything about cardiology, but aspires to look at an ECG graph and be able to detect *Atrial Fibrillation (AF)*. If this was her only goal and didn't care much about cardiology or medicine, what steps should Alice take in order to achieve her aspiration?

# Prerequisites for Intelligence

- Suppose Alice is a layperson who doesn't know anything about cardiology, but aspires to look at an ECG graph and be able to detect *Atrial Fibrillation (AF)*. If this was her only goal and didn't care much about cardiology or medicine, what steps should Alice take in order to achieve her aspiration?

# Prerequisites for Intelligence

- Suppose Alice is a layperson who doesn't know anything about cardiology, but aspires to look at an ECG graph and be able to detect *Atrial Fibrillation (AF)*. If this was her only goal and didn't care much about cardiology or medicine, what steps should Alice take in order to achieve her aspiration?
- ① Obtain two sets of ECGs, one set that was diagnosed as AF and the other with no AF.



# Prerequisites for Intelligence

- Suppose Alice is a layperson who doesn't know anything about cardiology, but aspires to look at an ECG graph and be able to detect *Atrial Fibrillation (AF)*. If this was her only goal and didn't care much about cardiology or medicine, what steps should Alice take in order to achieve her aspiration?
- - ① Obtain two sets of ECGs, one set that was diagnosed as AF and the other with no AF.
  - ② Collectively look at the two sets and look for a *pattern* that is only found in the set that has AF.

# Prerequisites for Intelligence

- Suppose Alice is a layperson who doesn't know anything about cardiology, but aspires to look at an ECG graph and be able to detect *Atrial Fibrillation (AF)*. If this was her only goal and didn't care much about cardiology or medicine, what steps should Alice take in order to achieve her aspiration?
- - ① Obtain two sets of ECGs, one set that was diagnosed as AF and the other with no AF.
  - ② Collectively look at the two sets and look for a *pattern* that is only found in the set that has AF.
  - ③ If no pattern was identified, repeat from step 1 (more data).



# Prerequisites for Intelligence

- Suppose Alice is a layperson who doesn't know anything about cardiology, but aspires to look at an ECG graph and be able to detect *Atrial Fibrillation (AF)*. If this was her only goal and didn't care much about cardiology or medicine, what steps should Alice take in order to achieve her aspiration?
- - ① Obtain two sets of ECGs, one set that was diagnosed as AF and the other with no AF.
  - ② Collectively look at the two sets and look for a *pattern* that is only found in the set that has AF.
  - ③ If no pattern was identified, repeat from step 1 (more data).
  - ④ If a pattern was identified, acquire a new undiagnosed ECG. If the identified pattern is present in the new ECG, detect AF, if not, no AF. Aspiration achieved!

# Prerequisites for Intelligence

- Suppose Alice is a layperson who doesn't know anything about cardiology, but aspires to look at an ECG graph and be able to detect *Atrial Fibrillation (AF)*. If this was her only goal and didn't care much about cardiology or medicine, what steps should Alice take in order to achieve her aspiration?
- - ① Obtain two sets of ECGs, one set that was diagnosed as AF and the other with no AF.
  - ② Collectively look at the two sets and look for a *pattern* that is only found in the set that has AF.
  - ③ If no pattern was identified, repeat from step 1 (more data).
  - ④ If a pattern was identified, acquire a new undiagnosed ECG. If the identified pattern is present in the new ECG, detect AF, if not, no AF. Aspiration achieved!
- The key idea here is that in order to possess the intelligence of being able to discern between ECG's, first Alice had to extract patterns from data.

# Prerequisites for Intelligence

- Suppose Alice is a layperson who doesn't know anything about cardiology, but aspires to look at an ECG graph and be able to detect *Atrial Fibrillation (AF)*. If this was her only goal and didn't care much about cardiology or medicine, what steps should Alice take in order to achieve her aspiration?
- - ① Obtain two sets of ECGs, one set that was diagnosed as AF and the other with no AF.
  - ② Collectively look at the two sets and look for a *pattern* that is only found in the set that has AF.
  - ③ If no pattern was identified, repeat from step 1 (more data).
  - ④ If a pattern was identified, acquire a new undiagnosed ECG. If the identified pattern is present in the new ECG, detect AF, if not, no AF. Aspiration achieved!
- The key idea here is that in order to possess the intelligence of being able to discern between ECG's, first Alice had to extract patterns from data.
- Therefore the ability to extract patterns from data is a crucial prerequisite for intelligence.

# What is Machine Learning?

- Machine Learning is the ability of a computer program to acquire knowledge by recognising hidden patterns in data.
- Key consideration: A pattern can be extracted only if there exists one.
- In this module we'll look at learning algorithms that can extract patterns out of data.



# Table of Contents

## 1 Module Outline

## 2 Introduction

- What is Intelligence?
- What is Machine Learning?

## 3 Learning Algorithms

## 4 Reference



# What is a Learning Algorithm?



# What is a Learning Algorithm?

- Informally:  
An algorithms that can detect patterns in data that is relevant to some task is called a learning algorithm.

# What is a Learning Algorithm?

- Informally:

An algorithm that can detect patterns in data that is relevant to some task is called a learning algorithm.

- More formally:

*A computer program is said to learn from experience  $E$  with respect to some task  $T$  and some performance measure  $P$ , if its performance on  $T$ , as measured by  $P$ , improves with experience  $E$*

- Tom Mitchell, 1997



# Components of a Learning Algorithm

- **Dataset:** Referred to as experience  $E$ , is a collection of example of performing the task  $T$ .

# Components of a Learning Algorithm

- **Dataset:** Referred to as experience  $E$ , is a collection of example of performing the task  $T$ .
- **Model:** An adaptable mathematical object that is used to explain the data in relation to the Task  $T$ . "Explaining the data" means the model's ability to represent accurately the underlying data-generating process.

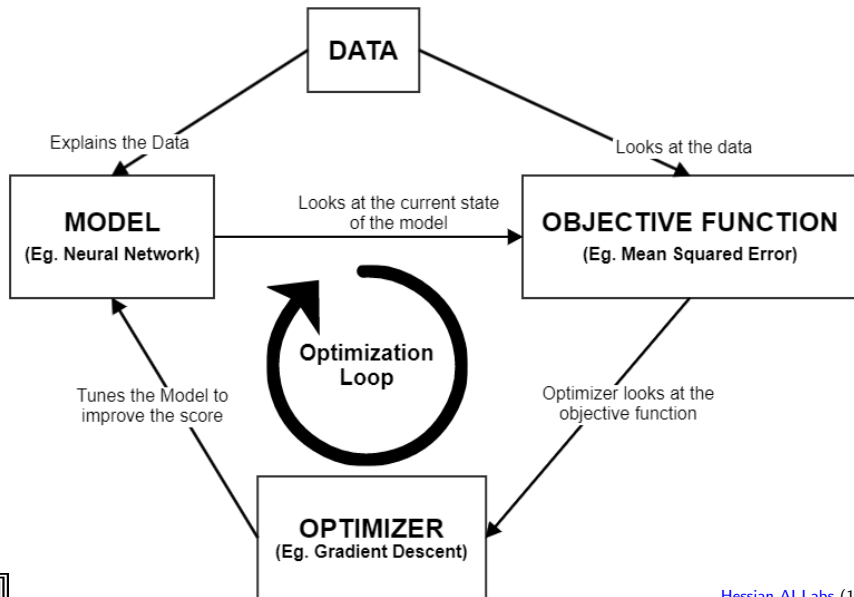
# Components of a Learning Algorithm

- **Dataset:** Referred to as experience  $E$ , is a collection of example of performing the task  $T$ .
- **Model:** An adaptable mathematical object that is used to explain the data in relation to the Task  $T$ . "Explaining the data" means the model's ability to represent accurately the underlying data-generating process.
- **Objective Function:** A function that scores the current state of the model based on its ability to accurately explain the data in relation to task  $T$ .

# Components of a Learning Algorithm

- **Dataset:** Referred to as experience  $E$ , is a collection of example of performing the task  $T$ .
- **Model:** An adaptable mathematical object that is used to explain the data in relation to the Task  $T$ . "Explaining the data" means the model's ability to represent accurately the underlying data-generating process.
- **Objective Function:** A function that scores the current state of the model based on its ability to accurately explain the data in relation to task  $T$ .
- **Optimizer:** The optimizer tunes the adaptable parameters of the model such that it improves the model's ability to explain the data, i.e. improves the score. When the objective function is a loss function the optimizer minimizes the score, when it is a performance measure the optimizer maximizes the score.

# Simplified Learning Loop



# Homework

- Finish Lab1 and Lab2 by Sunday April 19th.
- Read pages 1-3 (till before section 1.1) from PRML Bishop Chapter 1

# Table of Contents

- 1 Module Outline
- 2 Introduction
  - What is Intelligence?
  - What is Machine Learning?
- 3 Learning Algorithms
- 4 Reference



- Universal Intelligence: A Definition of Machine Intelligence (Legg and Hutter, 2007): <https://arxiv.org/abs/0712.3329>
- Mitchell, T. M. (1997). Machine learning. 1997. Burr Ridge, IL: McGraw Hill, 45(37), 870-877.